

Lecture 5

Large-Scale Network Attacks

ENGG5105/CSCI5470 Computer and Network Security

Spring 2014

Patrick P. C. Lee

Review of Last Lecture

- Many network attacks exploit the weaknesses of the fundamental design features of network protocols

Topics to Cover

➤ Port scanning

- A technique that quickly searches for vulnerabilities on many hosts

➤ Worms

- large-scale scanning for compromised hosts

➤ DDoS attacks

- denial of service attacks from many compromised hosts

➤ Botnets

- An army of compromised hosts

Our Goals

- To understand existing attack strategies of large-scale attacks
- To identify novel attacks that are potentially possible
- To help us develop defense strategies for existing and new attacks

Roadmap

➤ Port Scanning

➤ DoS

➤ DDoS

➤ Worms

➤ Botnets

Port Scanning

- Each network process on a single host is associated with a **port number**, so that incoming packets can identify which network process they are destined for.
 - It's like a door.
- Port numbers ranging from 0 – 1023 are called **well-known port numbers** and reserved for use by well-known applications
 - Examples:
 - HTTP: TCP 80
 - FTP: 20 (data), 21 (control)
 - DNS: 53
- Check: <http://www.iana.org/assignments/port-numbers>

Port Scanning

- Port scanning is to scan through a range of ports to identify any active network process
 - **Horizontal portscans**: scan a range of hosts on a specific port
 - **Vertical portscans**: scan a range of ports on a specific host
- An attacker scans for open ports to find vulnerable network processes to compromise

Nmap

- Nmap is a full-featured port scanning tool
 - used by operators to debug network services
- Nmap provides different scan types, each sends a specific type of packets to a target.
- Some scan types could cause the target system to be flooded or even crash under the load of strange and unusual types

Scan Types in Nmap

➤ TCP connect() scan:

- completes the 3-way handshake with each scan port
- **not stealthy** (i.e., source IP easily gets logged when a connection is open)
- command line option: -sT

Scan Types in Nmap

➤ TCP SYN scan:

- only sends the initial SYN and awaits the SYN-ACK response to determine if a port is open. If the port is closed, the destination will send a RST or nothing.
- Stealthier than TCP connect() since only two packets are involved for open ports (while TCP connect() involves three packets)
- command line option: -sS

Scan Types in Nmap

➤ TCP FIN scan:

- sends a TCP FIN to each port. The FIN is unexpected as no connection is created before. A RESET indicates the port is closed, and no response may mean the port is open.
- may be broken for systems that send RESET anyway
- Even stealthier than TCP SYN (why?)
- command line option: -sF

Scan Types in Nmap

➤ OS Fingerprinting:

- Instead of scanning open ports, Nmap uses it to determine the OS type of the target based on their responses to illegal combinations of TCP control bits
- command line option: -O

Defenses Against Port Scanning

- Close all unused ports
 - Use Nmap:
 - `nmap localhost`
 - Check for listening ports:
 - `netstat -na | grep "LISTEN"`
- Apply firewalls to block port accesses
 - iptables
- Use network intrusion detection system
 - Not blocking scanners, but detecting the presence of scanners

Roadmap

➤ Port Scanning

➤ DoS

➤ DDoS

➤ Worms

➤ Botnets

Denial-of-Service (DoS)

- In the last lecture, we focus on network attacks that gain access
- **Denial-of-service (DoS)** attacks deny access of legitimate users or stop critical system processes
- Launching DoS attacks has a business incentive
 - e.g., you want your competitor companies on their e-commerce websites

Ways of Launching DoS Attacks

- DoS attacks can be launched locally or remotely.
 - **Locally**: install a local account on a compromised machine
 - e.g., kill/crash a process, reconfigure systems
 - **Remotely**: attackers generate attack traffic to a target machine across the network
- Remote DoS attack is more prevalent

Types of DoS Attacks

➤ Stopping a service

- Crash or shutting off a specific program or machine that users want to access

➤ Resource exhaustion

- While the service is still running, attackers consume computer or network resources to prevent users from reaching the service
 - e.g., by **flooding** a large amount of packets

➤ In this lecture, we focus on network-based (remote) DoS attacks based on resource exhaustion

SYN Flood

- All TCP connections begin with 3-way handshake, in which a client first sends a SYN packet to an open port of a server
- When the server receives SYN, it allocates state to track the status of the half-open connection, and replies SYN-ACK
- In **SYN flood**, an attacker sends a large number of SYN packets to a target server to consume its resources

SYN Flood

- Attack SYN packets usually have spoofed IP addresses that are not existing (unresponsive)
 - Forcing the server to keep the state for a long while
 - By choosing a large pool of spoofed addresses, the detection will be difficult
- Two ways of resource exhaustion
 - Fill up the connection queue of the server
 - Fill up the communication link attached to the server

SYN Flood Defenses

- Provision enough resources
 - May not be feasible
- SYN Cookies
 - only need to be implemented on the TCP/IP stack of the server
 - carefully construct sequence numbers in the SYN-ACK field sent by the server
 - trade CPU for memory

SYN Cookies

- How traditional 3-way handshake works?
 - $A \rightarrow B: \text{SYN}(A, \text{ISN}_A)$
 - $B \rightarrow A: \text{SYN}(B, \text{ISN}_B), \text{ACK}(A, \text{ISN}_A+1)$
 - $A \rightarrow B: \text{ACK}(B, \text{ISN}_B+1)$
- ISN_B is a cryptographic one-way function value of source/destination IP addresses, port numbers, time, and a secret seed. ISN_B is called a SYN cookie.
- When Bob replies SYN-ACK, he **forgets** the half-open state, i.e., no connection queue is used.

SYN Cookies

- When the $\text{ACK}(B, \text{ISN}_B + 1)$ arrives, Bob applies the same function to validate if ISN_B is legitimate. If so, the connection is established.
 - Need additional computation on hash
- Trudy sends spoofed SYN packets, but Bob doesn't store any state.
- Of course, Trudy may send a flood of ACKs to consume CPU resources of Bob. That's the tradeoff.


SYN Cookies

- In Linux, you can enable SYN cookies:
 - `echo "1" > /proc/sys/net/ipv4/tcp_syncookies`
- Other defenses against SYN flooding
 - **traffic shaping**: apply rate-limiting on routers to upper-bound the number of incoming SYN packets

Smurf Attacks

- Smurf attacks use direct broadcast to create a flood of traffic for a victim
- How broadcast works?
 - If a packet has IP address *.*.*.255, then it's a broadcast packet
 - Directed broadcast: host part of an address is all 1's (e.g., 10.1.255.255 for 10.1.0.0/16)
 - When a switch/router receives a broadcast packet for a LAN, it sends a packet to every system on the LAN

Smurf Attacks



- An attacker sends a ping request to a broadcast address (ICMP Echo Request)
- Every host on the target network makes a replies to a victim.
- The victim is overloaded with many ping replies.
- **Defense:** disable directed broadcasts in routers

Roadmap

- Port Scanning
- DoS
- DDoS
- Worms
- Botnets

Distributed DoS

- DoS attacks usually refer to point-to-point attacks:
 - one attack source, one target
- **Distributed DoS (DDoS)** attacks usually refer to multipoint-to-point attacks:
 - an attacker controls multiple attack sources to attack a target
- But you need to control many attack sources first before launching DDoS attacks.
 - How???

DDoS

- DDoS is difficult to detect
 - There are many attack sources, so finding all of them is difficult
 - Each attack source only needs to send a small amount of attack traffic, but the aggregate can be overwhelming


DDoS Flood Types

➤ Direct DoS attacks:


- instruct zombies to send flood traffic

➤ Reflector DoS attacks

- Send request packets to uninfected machines, while packets have spoofed source IP address set to the target
- All uninfected machines make replies to the target
- e.g., send DNS queries to DNS servers



(a) Direct DDoS Attack



(b) Reflector DDoS Attack

DDoS Countermeasures

- three broad lines of defense:
 1. attack prevention & preemption (before)
 2. attack detection & filtering (during)
 3. attack source traceback & ident (after)
- huge range of attack possibilities
- hence evolving countermeasures

Roadmap

- Port Scanning
- DoS
- DDoS
- Worms
- Botnets

Viruses

- piece of software that infects programs
 - modifying them to include a copy of the virus
 - so it executes secretly when host program is run
- specific to operating system and hardware
 - taking advantage of their details and weaknesses

Worms


- How do you control many attack sources first before launching DDoS attacks?
- **Worms** are programs that **self-propagate** across the Internet by exploiting security flaws in widely-used services
 - E.g., by port scanning
- Worms are different from viruses, where the latter require some sort of user actions to abet their propagation
 - but some worms also requires user to trigger, but their propagation are done by themselves
 - won't cover viruses in details in this course.

Morris Worm

- one of best known worms
- released by Robert Morris in 1988
- various attacks on UNIX systems
 - cracking password file to use login/password to logon to other systems
 - exploiting a bug in the finger protocol
 - exploiting a bug in sendmail
- if succeed have remote shell access
 - sent bootstrap program to copy worm over

How Serious are Worms?

From [Staniford et al., 2002]



- Onset of Code Red I v2, code Red II and Nimda: Number of remote hosts launching confirmed attacks corresponding to different worms, as seen at the Lawrence Berkeley National Laboratory.

Code Red I

- Launched on July 13, 2001 (version 2 launched on July 19, 2001).
- Once Code Red I infected a host, it spreads by launching 99 threads that generated random IP addresses.
- v2 launches payload to www.whitehouse.gov

Worm Propagation Model

- Propagation model of Code Red I is based on **Random Constant Spread (RCS)**
- How to provide a mathematical model for it?
- Let
 - **N** = number of vulnerable hosts in the Internet
 - **K** = compromise rate: number of vulnerable hosts which an infected host can find and compromise per unit time
 - **a** = proportion of vulnerable machines that have been compromised

Worm Propagation Model

➤ We can construct a differential equation:

$$N da = (Na) K(1 - a)dt$$

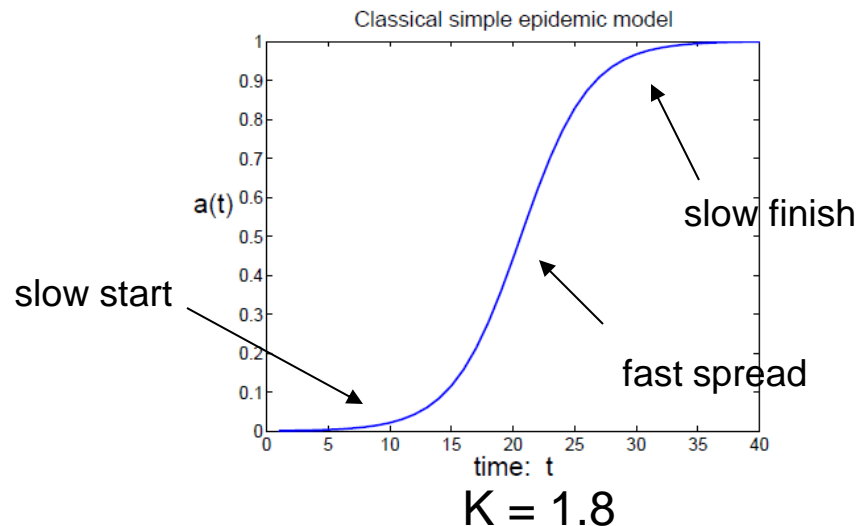
Number of machines compromised in the next increment of time ($N da$) is equal to the number of machines already compromised (Na) times the number of machines each compromised machine can compromise ($K(1-a)dt$).

Worm Propagation Model

➤ $\frac{da}{dt} = Ka(1-a)$

➤ Three phases:

- Slow-start phase:
 - if a is small, $\frac{da}{dt} \sim Ka$ (exponential growth)
- Fast-spread phase:
 - $\frac{da}{dt}$ is max when $a=0.5$
- Slow-finish phase:
 - if $a \rightarrow 1$, $\frac{da}{dt} \sim 0$ (all vulnerable machines are compromised and no leftover)



From “Code Red Worm Propagation Modeling and Analysis”, CCS 2002

Code Red II – Better Worm

- Code Red I chose a random IP address from the Internet and attempted to infect it
- Code Red II (released on Aug 4, 2001) used a localized scanning strategy, which scans the addresses that close to it
 - probability $3/8$ to choose an IP from the $/16$ network of the infected machine itself
 - probability $1/2$ to choose an IP from its own $/8$
 - probability $1/8$ to choose a random address from the Internet

Nimda – Better Worm

- Nimda (released on Sep 18, 2001) used different methods for its propagation
 - probing for a Microsoft IIS vulnerability
 - bulk emailing of itself as an attachment
 - copying itself across open network shares
 - adding exploit code to Web pages
 - scanning the backdoors left by Code Red II

Recent Worm Attacks

➤ SQL Slammer

- early 2003, attacks MS SQL Server

➤ Mydoom

- mass-mailing e-mail worm that appeared in 2004
- installed remote access backdoor in infected systems

➤ Warezov family of worms

- scan for e-mail addresses, send in attachment

Mobile Phone Worms

- first appeared on mobile phones in 2004
 - target smartphone
- they communicate via Bluetooth or MMS
- to disable phone, delete data on phone, or send premium-priced messages
- CommWarrior, launched in 2005
 - replicates using Bluetooth to nearby phones
 - and via MMS using address-book numbers

Worm Countermeasures

- overlaps with anti-virus techniques
- once worm on system A/V can detect
- worms also cause significant net activity
- worm defense approaches include:
 - signature-based worm scan filtering
 - filter-based worm containment
 - payload-classification-based worm containment
 - threshold random walk scan detection
 - rate limiting and rate halting

Roadmap

- Port Scanning
- DoS
- DDoS
- Worms
- Botnets

Botnets


- Using worms, an attacker can compromise millions of hosts, and form a **botnet**
- A **bot** is a compromised host that can be remotely controlled by an attacker (called **botmaster**)
- A **botnet** is a network of bots (or bot army)

Centralized Botnets

- Most existing botnets use a centralized architecture
- A botmaster sends commands to **Command and Control (C&C)** servers, which relay the commands to bots
 - e.g., the command could be “send junk packets to info.gov.hk”
 - using Internet Relay Chat (IRC) protocol – send commands, receive replies
- Weaknesses:
 - Easier to detect
 - Single point of failure (from an attacker’s view)

Centralized Botnets

➤ Centralized C&C-based Botnet



A defender can fix one of the C&C servers and find out where the botmaster is

Decentralized Botnets

- **Peer-to-peer (P2P):** a botmaster sends commands to its peer bot through an existing P2P network
 - More resilient connectivity
 - Difficult to be detected
- **Unstructured:** pass along encrypted commands to a random peer (may or may not be bot). If the peer is a bot, it decrypts and executes the commands.

Attacks and Theft

➤ Why creating a botnet?

- Launch DDoS attacks to enemies
- Collect personal information (e.g., credit card numbers) and make business
- Send spam (most important use)
- Create phishing websites

Honeypots

- A **honeypot** is a decoy system to lure attackers
 - unpatched system so as to be attacked
 - away from accessing critical systems
 - to collect information of attackers' activities
 - to encourage attackers to stay on system so administrator can respond
- Usually deployed as a virtual machine
- A **honeynet** is a network of honeypots

References

- “Counter Hack Reloaded”, Ch. 7, 9
- Stallings, Ch. 21
- Staniford et al., “How to Own the Internet in your Spare Time”, Security '02
- Bailey et al., “A Survey of Botnet Technology and Defenses”, 2009