

Lecture 6

Firewall and NIDS

ENGG5105/CSCI5470 Computer and Network Security


Spring 2014

Patrick P. C. Lee

How to Defend Against Network Attacks?


➤ **Prevention:** block/deny attacks

- e.g., firewall



➤ **Detection:** detect attacks, no immediate reaction

- e.g., NIDS (network intrusion detection system)



Roadmap


- Firewall
- NIDS (Network Intrusion Detection System)

What is a Firewall

- A **choke point** of control and monitoring
- Interconnects networks with differing trust
- Imposes restrictions on network services
 - only authorized traffic is allowed
- Auditing and controlling access
 - can implement alarms for abnormal behavior
- Must be immune to penetration

Deploying Firewalls

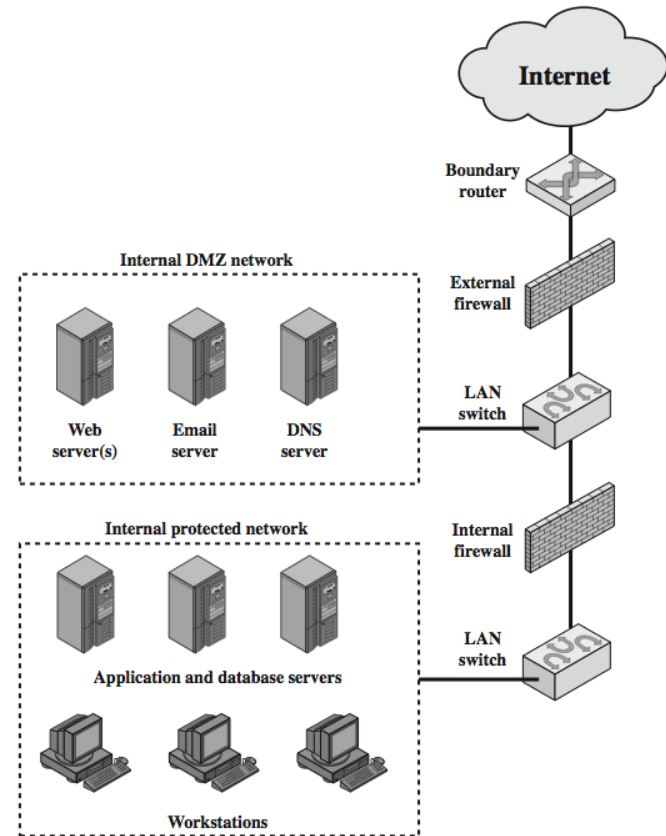
- We can deploy a firewall that interconnects external network (untrusted) and internal network (trusted)



(a) Packet-filtering router

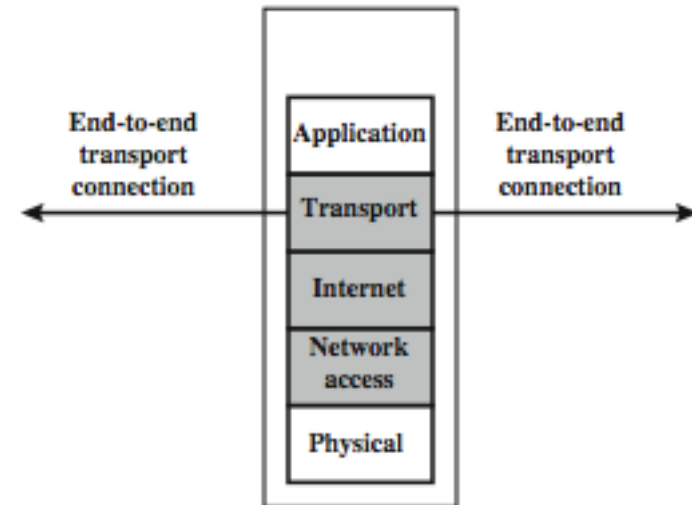
Deploying Firewalls

- Deploy multiple firewalls with different policies
- Internal firewalls have more stringent policies to protect internal hosts
- A **DMZ (demilitarized zone)** between the external and internal firewalls to deploy public services



Firewall – Packet Filters

- Simplest, fastest firewall component
- Foundation of any firewall system
- Examine each IP packet and permit or deny according to **rules**
 - e.g., you can define a rule that blocks any inbound packets to port 53 (DNS service)
 - hence restrict access to services (ports)



Firewalls – Packet Filters

Table 20.1 Packet-Filtering Examples

A	action	ourhost	port	theirhost	port	comment
	block	*	*	SPIGOT	*	we don't trust these people
	allow	OUR-GW	25	*	*	connection to our SMTP port

B	action	ourhost	port	theirhost	port	comment
	block	*	*	*	*	default

C	action	ourhost	port	theirhost	port	comment
	allow	*	*	*	25	connection to their SMTP port

D	action	src	port	dest	port	flags	comment
	allow	{our hosts}	*	*	25		our packets to their SMTP port
	allow	*	25	*	*	ACK	their replies

E	action	src	port	dest	port	flags	comment
	allow	{our hosts}	*	*	*		our outgoing calls
	allow	*	*	*	*	ACK	replies to our calls
	allow	*	*	*	>1024		traffic to nonservers

Firewall – Packet Filters

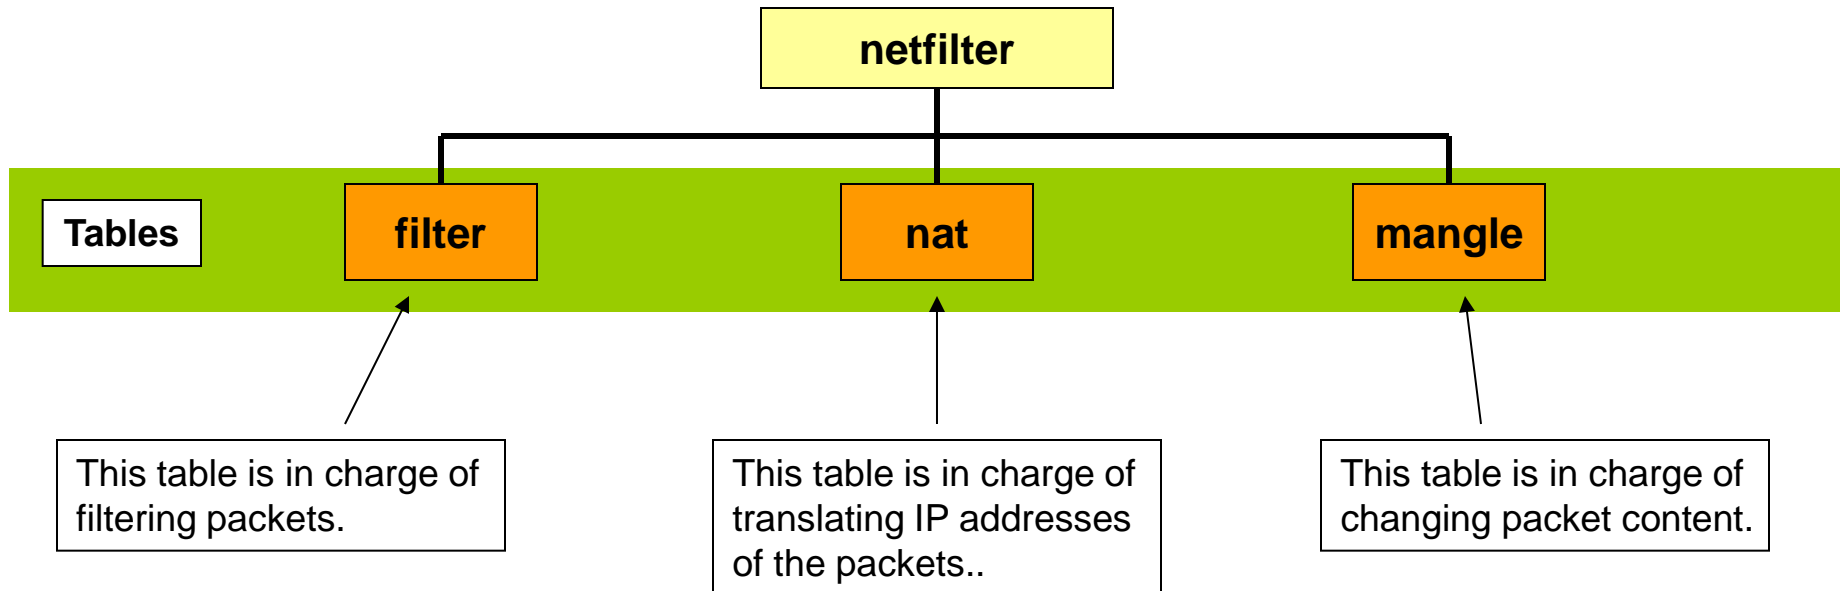
- Can implement **stateful filtering**
- Stateful packet filters examine each IP packet **in context**
 - keep track of client-server sessions
 - check each packet validly belongs to one
- Add more intelligence, but with higher processing overhead

iptables

- **iptables** is a user-level Linux program that controls the kernel-level network module called netfilter. It can perform:
 - Packet Filtering
 - Packet Forwarding
 - Network Address Translation (NAT)
 - Connection Tracking (stateful tracking)
- We can use iptables to implement a filter-based firewall atop Linux.

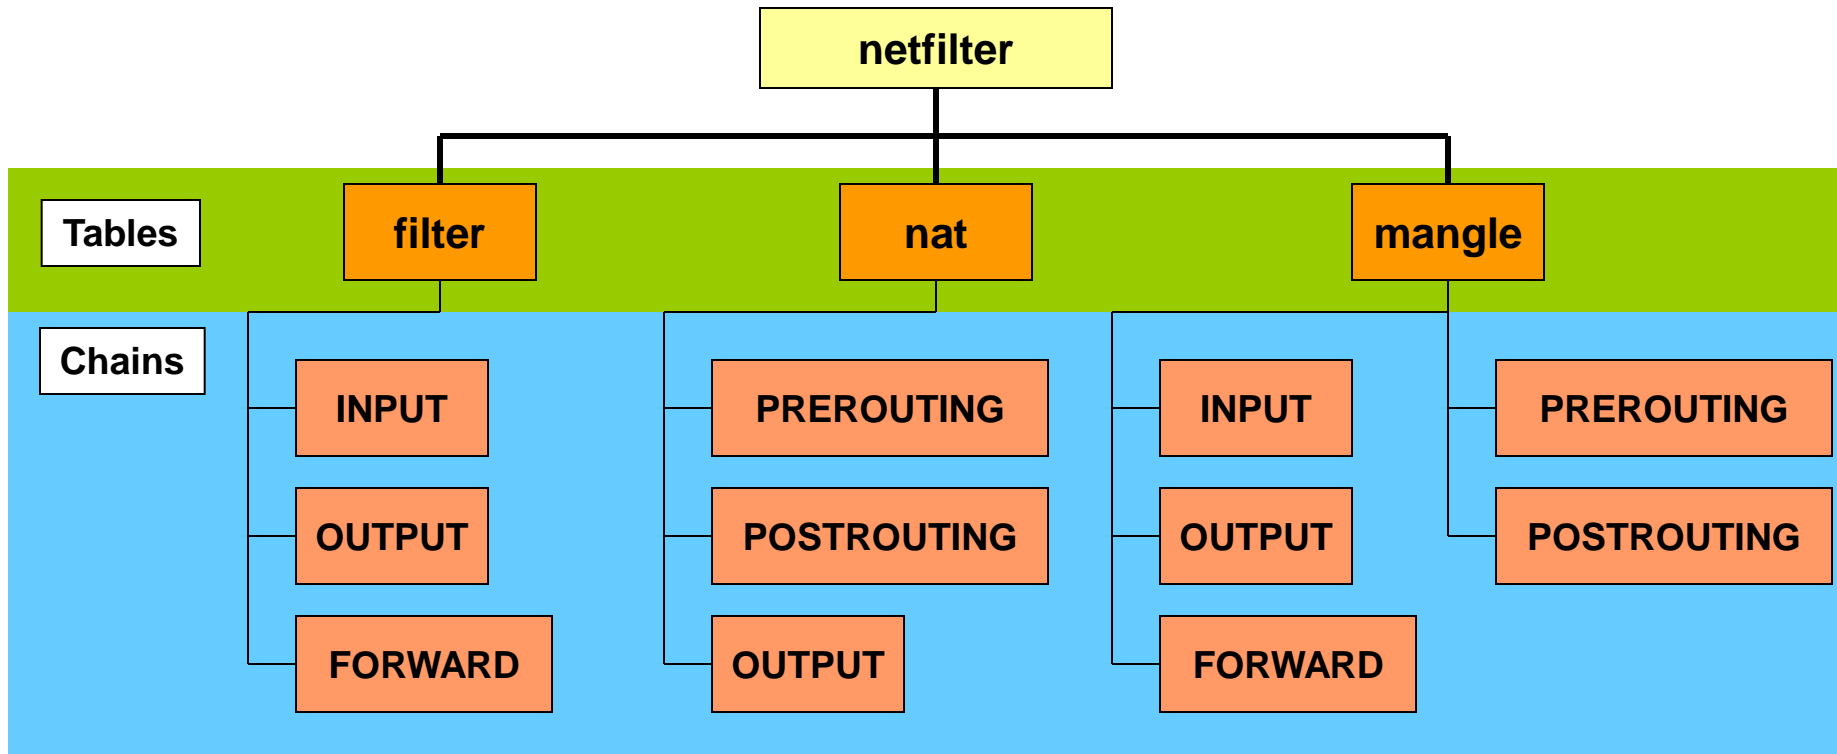
iptables – Tables and Chains

- Each function provided by the netfilter architecture is presented as a **table**.



iptables – Tables and Chains

- Under each table, there are a set of **chains**.
 - Under each chain, you can assign a set of **rules**.



iptables – Tables and Chains

Chain name: **INPUT**

Table name: **filter**

The command: **list**

There is one rule set in the INPUT chain.

The other two chains.

```
[root@linux]# iptables -t filter -L
Chain INPUT (policy ACCEPT)
target     prot opt source                               destination
DROP       icmp -- anywhere                          anywhere

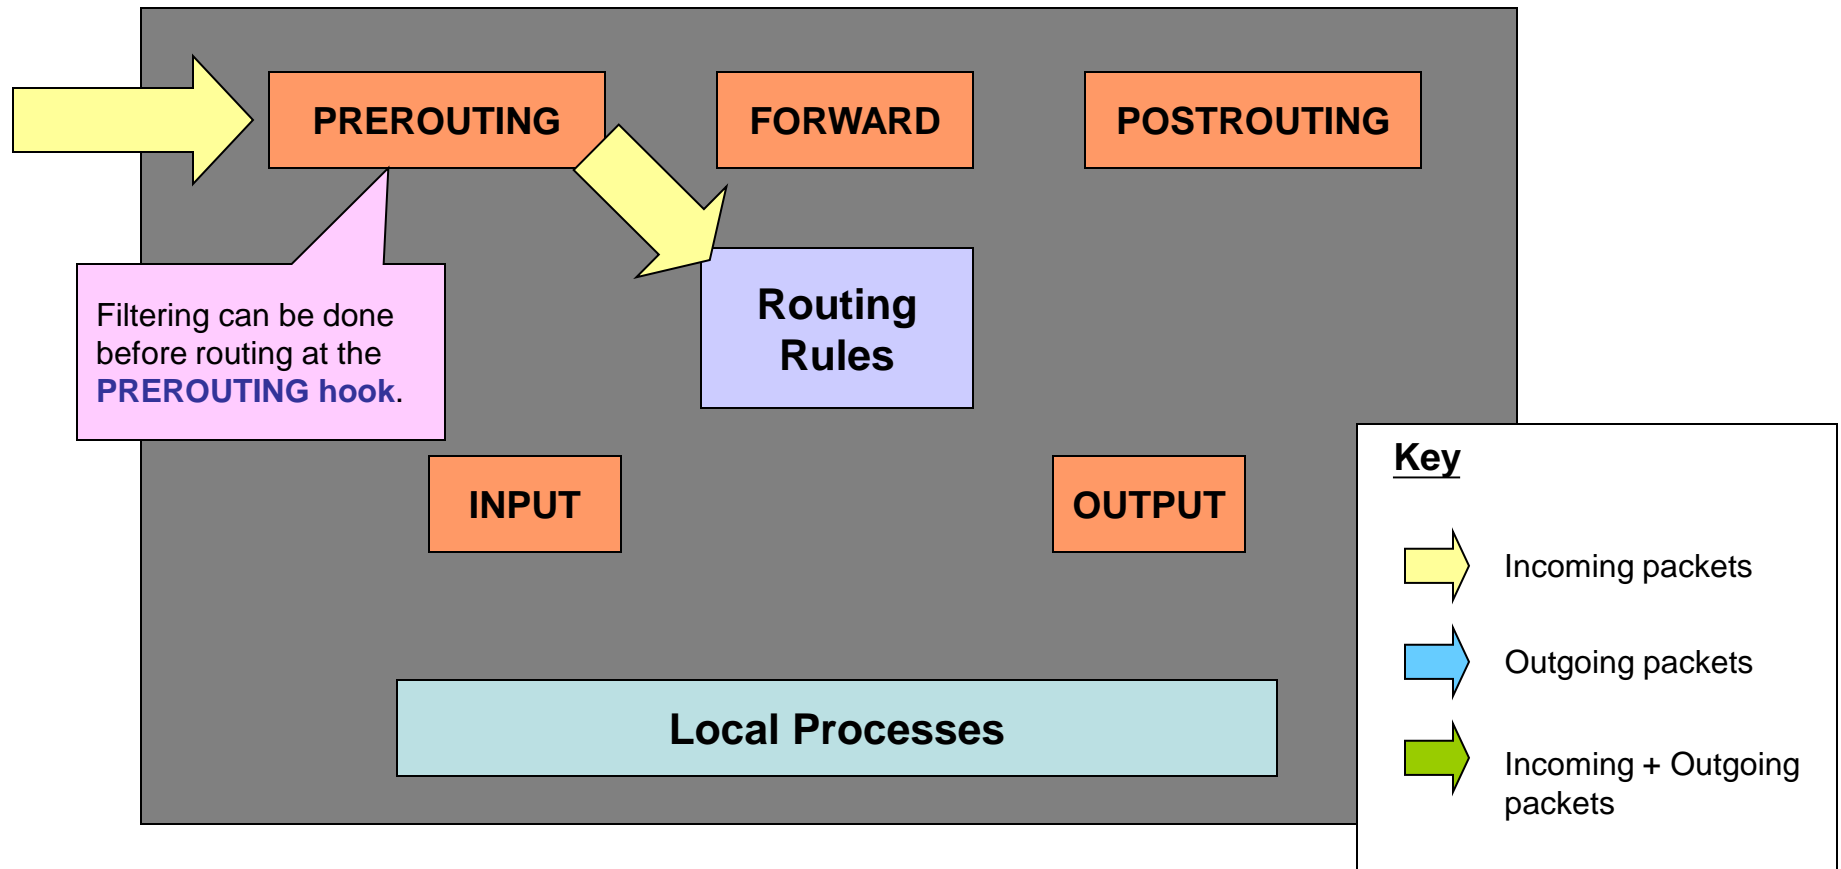
Chain FORWARD (policy ACCEPT)
target     prot opt source                               destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                               destination
[root@linux]# _
```

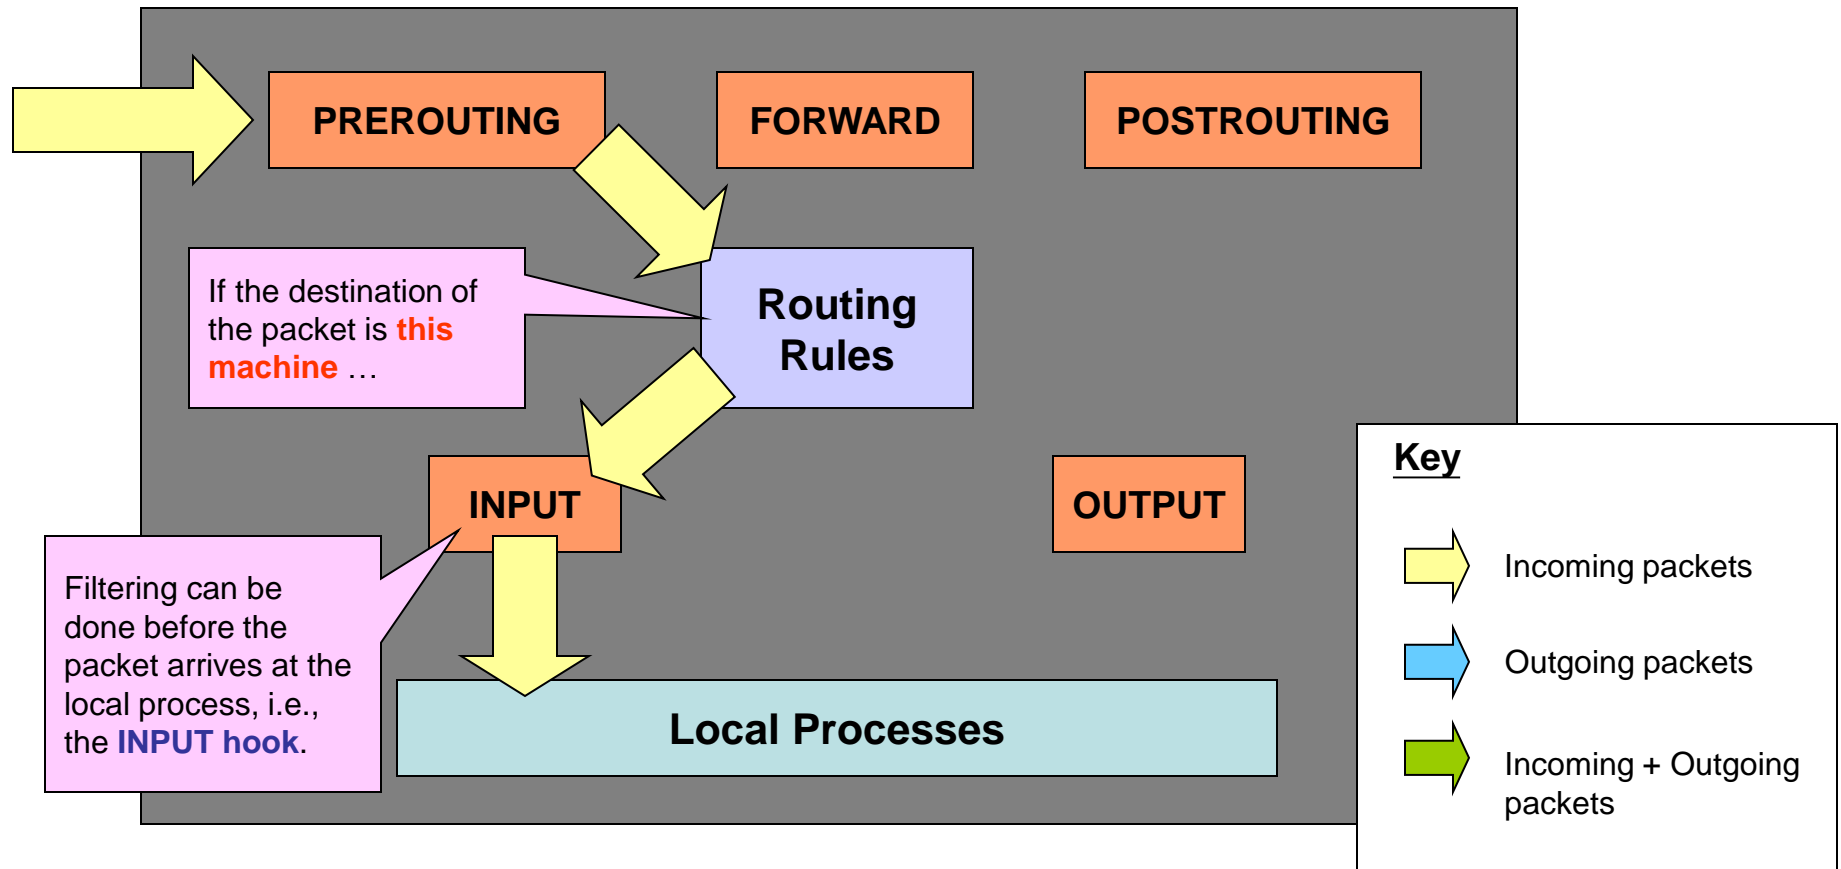
The rule in the INPUT chain means:

When a packet with ICMP payload passes through the **INPUT hook**, **DROP** that packets, no matter it is **from anywhere** and **to anywhere**.

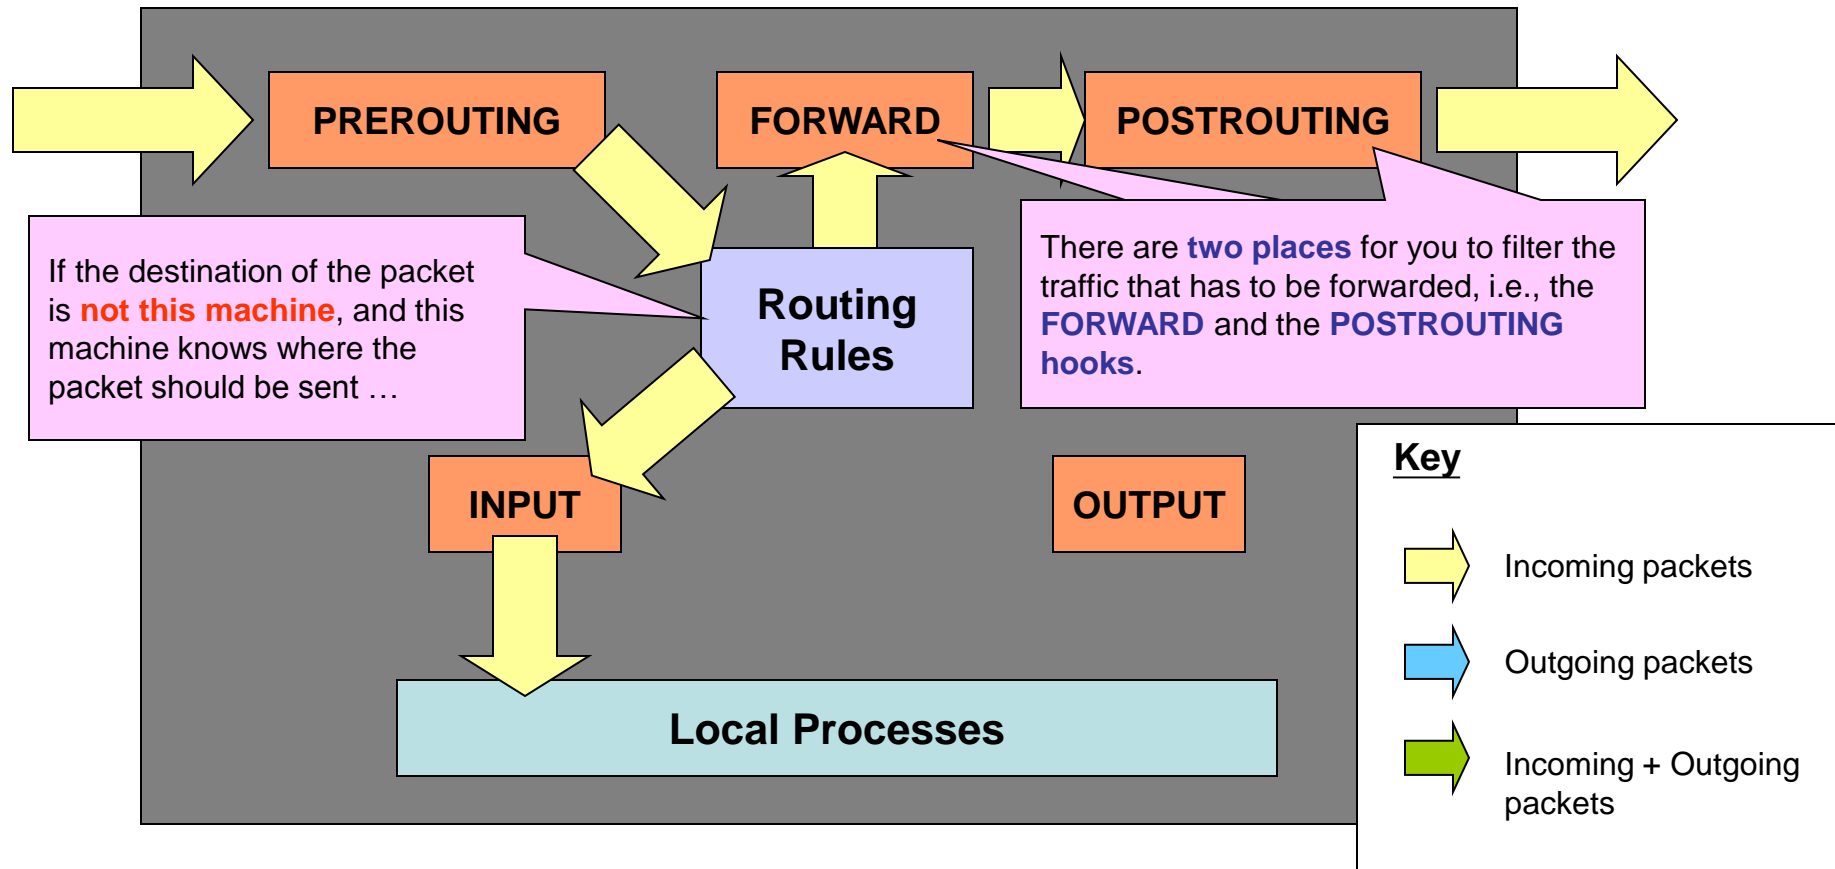
iptables – Packet Flow



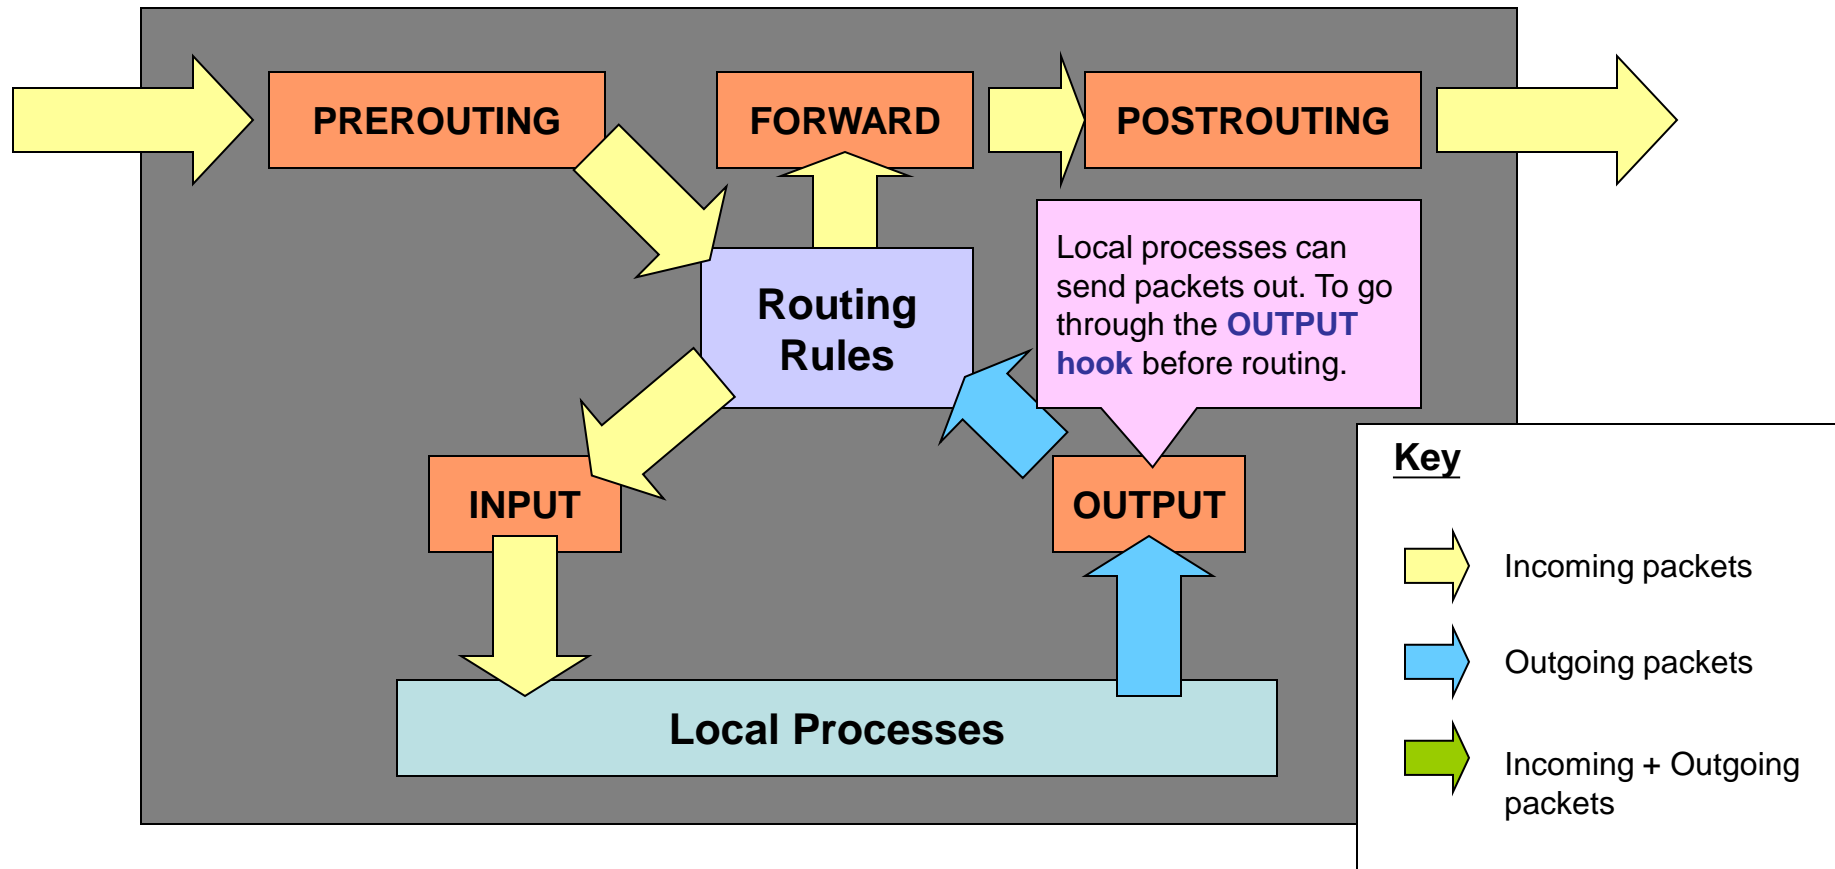
iptables – Packet Flow



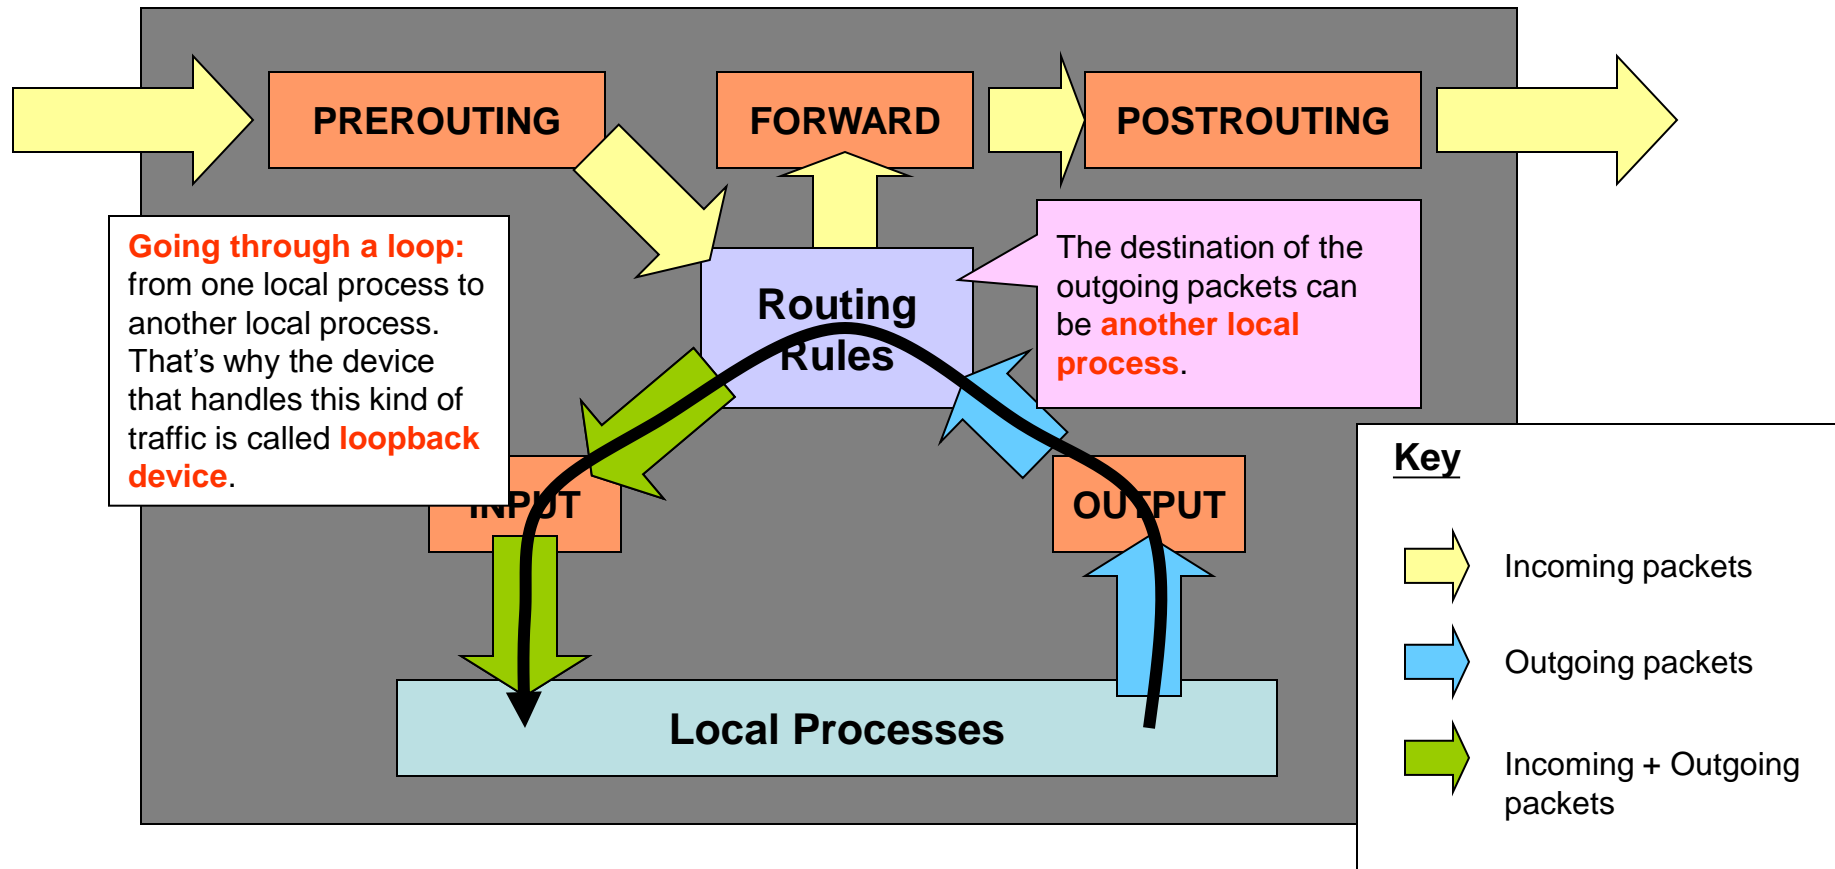
iptables – Packet Flow



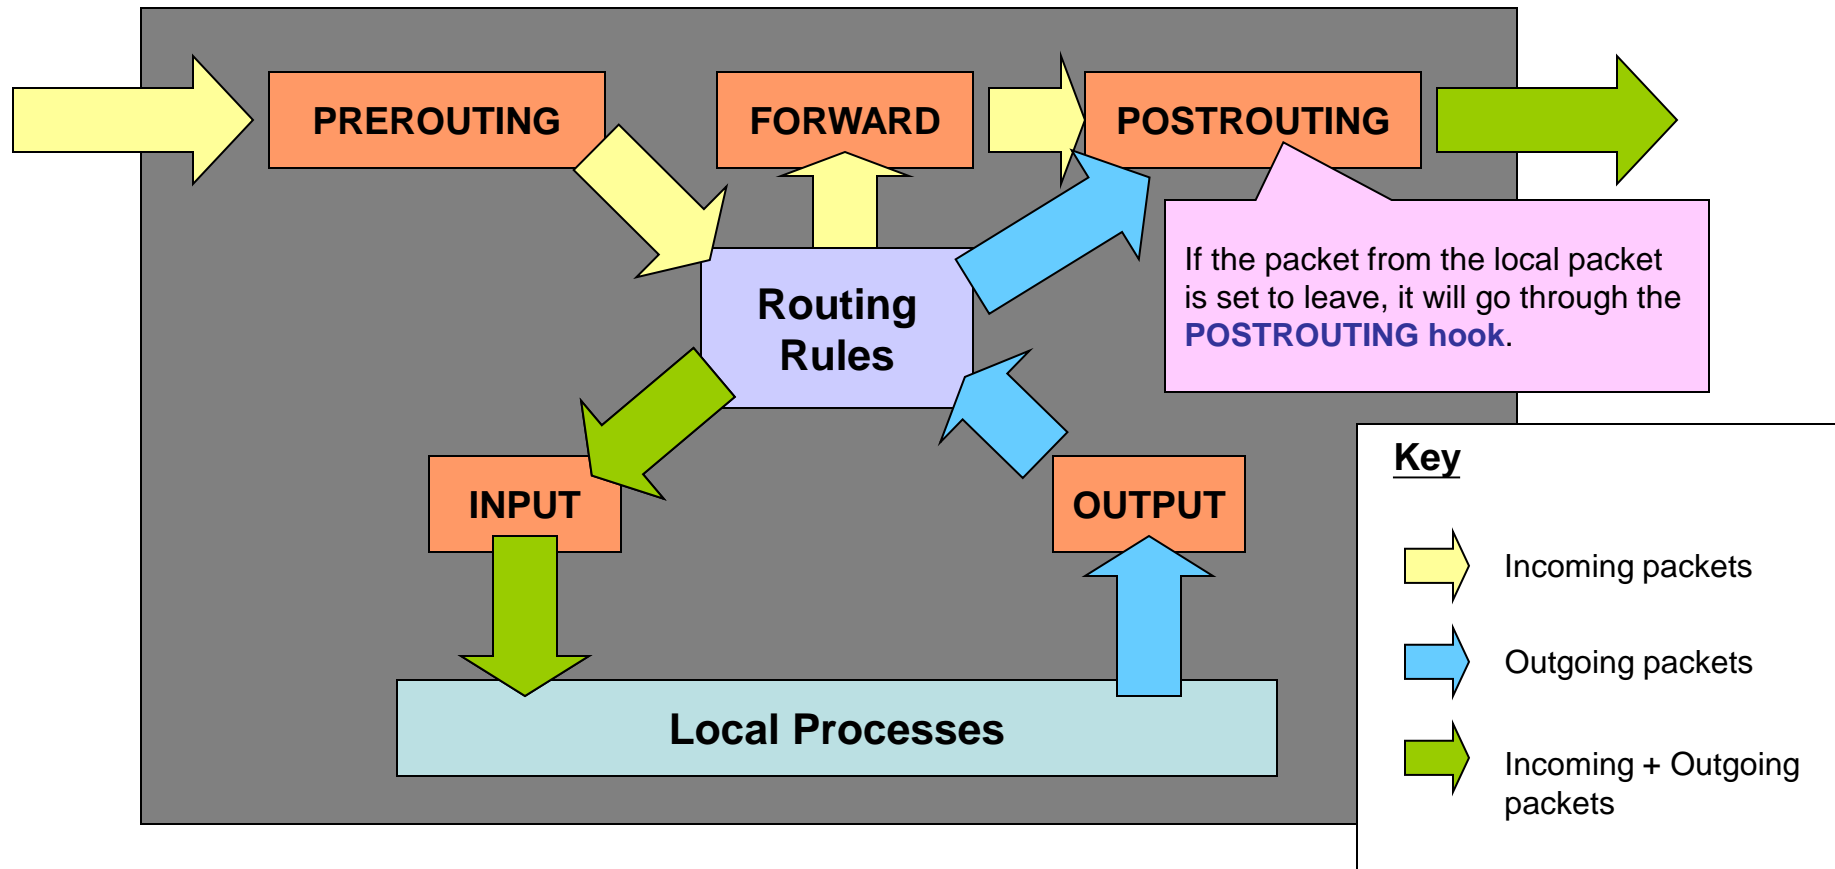
iptables – Packet Flow



iptables – Packet Flow



iptables – Packet Flow



Examples of Using iptables

- List all existing rules (default is filter table):

```
% iptables -L
```

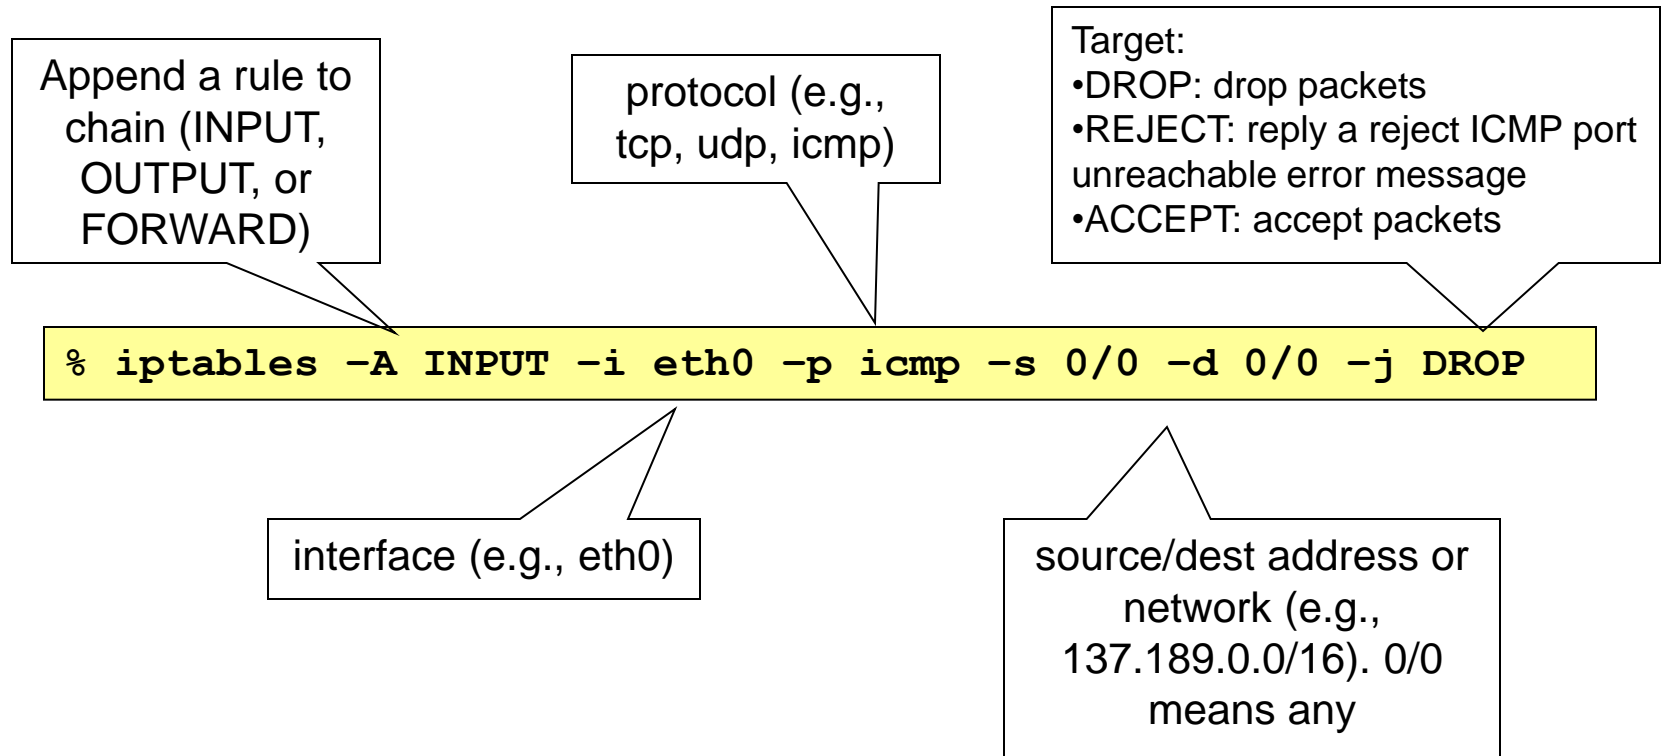
- Flush all existing rules:

```
% iptables -F
```

- Drop all incoming packets to eth0:

```
% iptables -A INPUT -i eth0 -p icmp -s 0/0 -d 0/0 -j DROP
```

Examples of Using iptables



Stateless Packet Filtering

- Accept the web service from CSE machines

```
% iptables -A INPUT -p tcp -s 137.189.88.0/22 -d 0/0 --dport 80 -j ACCEPT
```

- Reject all incoming TCP traffic with flags SYN, SYNACK, or RST that is destined for ports 0 to 1023

```
% iptables -A INPUT -p tcp -s 0/0 -d 0/0 --dport 0:1023 --syn -j REJECT
```

- Reject all outgoing TCP traffic except the one destined for 137.189.5.7

```
% iptables -A OUTPUT -p tcp -s 0/0 -d ! 137.189.5.7 -j REJECT
```

Stateful Packet Filtering

- Stateful inspection with option `-m`
- Use “`-m limit`” for SYN flood protection (limit rate to 1/s)

```
% iptables -A FORWARD -p tcp --syn -m limit --limit 1/s -j ACCEPT
```

- Use “`-m state`” to drop SYN/SYNACK/RST packets that will NOT create new connections

```
% iptables -A INPUT -p tcp --syn -m state ! --state NEW -j DROP
```

Attacks on Firewall

- IP address spoofing
 - fake source address to be trusted
- source routing attacks
 - attacker sets a route other than default to bypass the firewall
- tiny fragment attacks
 - split header info over several tiny packets
 - if the firewall doesn't do packet reassembly, its rules will become useless

Limitations of Firewall

- Internal attacks cannot be protected
- Single point of failures
- Performance bottleneck

Roadmap

- Firewall
- NIDS (Network Intrusion Detection System)

NIDS

- NIDS is to passively monitor malicious events in network traffic
- If a bad event is identified, alert the operators, or log the event.
- It's passive, no denying bad requests
 - won't affect existing traffic

Types of Intrusion Detection

➤ signature-based

- Looking for attack patterns
- e.g., at least K scans in T seconds, or payload has string “0x02030441”
- need to frequently update attack signatures

➤ anomaly-based

- finding abnormal behavior from “normal” behavior from profiling
- usually based on statistical techniques
- can have high false alarms

Design Goals of NIDS

See [Paxson, 98]

- High-speed, large volume monitoring
 - packet processing speed should match link capacity
- No packet filter drops
 - should have a large buffer to handle traffic spikes
- Real-time notification
 - An attack should be located ASAP
- Mechanism separate from policy
 - policies can be flexibly changed without re-implementing the system

Design Goals of NIDS

- Extensible
 - Upgrade should be easy
- Avoid simple mistakes
 - policy definitions should be clear
- The monitor will be attacked
 - Our design should prepare that an attacker can know everything about the logic of our system
- Examples of NIDS: Bro, Snort

Snort

- **Snort** is an open-source libpcap-based packet sniffer and logger for lightweight network intrusion detection
 - <http://www.snort.org>
- Developed by Marty Roesch in 1998
- Design features of Snort
 - lightweight (cross platform, small source)
 - rule-based detection
 - stateful packet analysis
 - extensible

Getting Started with Snort

- Download tarball from <http://www.snort.org>
 - Need to have PCRE (Perl Compatible Regular Expressions) installed
 - `apt-get install libpcre3 libpcre3-dev`
 - The version that we use is 2.8.6.1.
- Installation is straightforward:
 - Untar the tarball
 - change to the snort source directory
 - `./configure, ./make`

What can Snort do?

- **Sniffer mode**: reads packets off the network and displays them on the console
- **Packet logger mode**: logs packets to disks
- **Network intrusion detection system mode**: analyze packets for matches against user-defined rules and perform actions
- **Inline mode**: obtains packets from iptables and causes iptables to take actions

Snort as a Sniffer

- Making Snort as a sniffer:

```
% snort -v -i eth0
```

- This enables the verbose mode of Snort. Snort will listen to interface eth0. This only prints IP and TCP/UDP/ICMP headers
- You can see the application data with -d

```
% snort -vd -i eth0
```

Snort as a Packet Logger

- You can record packets into disk with the “-l” option (e.g., in directory log/):

```
% snort -v -i eth0 -l ./log
```

- The log file is in pcap format. You can read the log file with “-r” option:

```
% snort -v -r log_file
```

Snort as an NIDS

- You can apply rules to decide what actions to be made to each packet, online or offline.

```
% snort -d -i eth0 -l ./log -c snort.conf      # read from eth0 online  
% snort -d -r pkt.pcap -l ./log -c snort.conf # read from file offline
```

- **snort.conf** is the name of the **rule file**
- In online intrusion detection, the **-v** switch should be left off for speed.

Snort Alerts

- Alerts are mainly used to log events of interest in NIDS
- When you run snort, include `-A` switch
- Common options:
 - `-A fast`: fast alert mode. Write alerts in simple format
 - `-A full`: full alert mode (default)
 - `-A none`: turn off alert mode

Writing Snort Rules

- Snort rules are in single line (add \ to the end of line if a rule spans multiple lines).

```
alert tcp $BAD any -> $GOOD any (flags: SF; msg: "SYN-FIN Scan";)
```

Rule Header

Rule Header

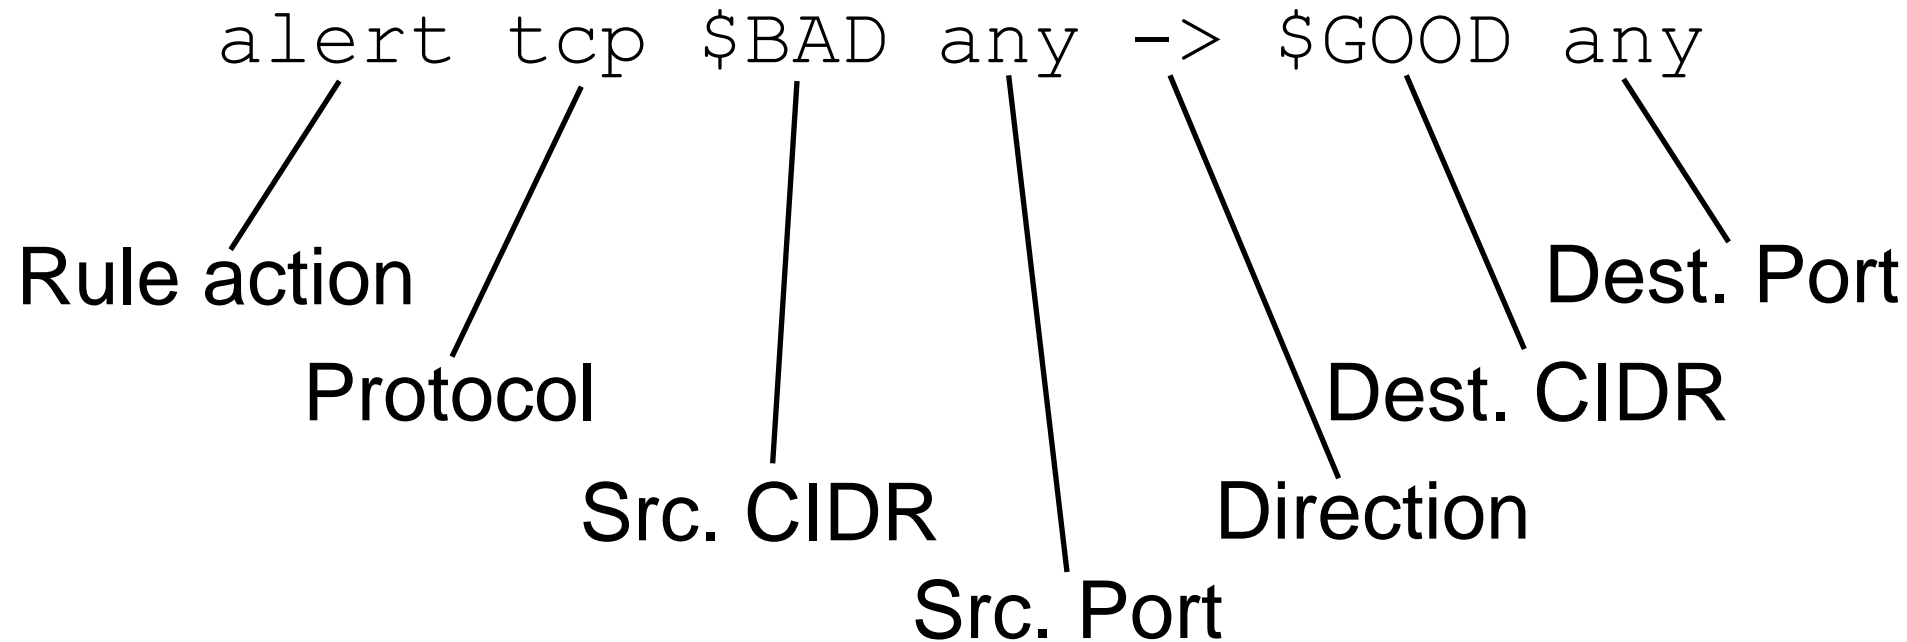
- static definition
- has to be in every rule

Rule Options

Rule Options

- variable definition
- not always necessary

Snort Rule Headers



Snort Rule Options

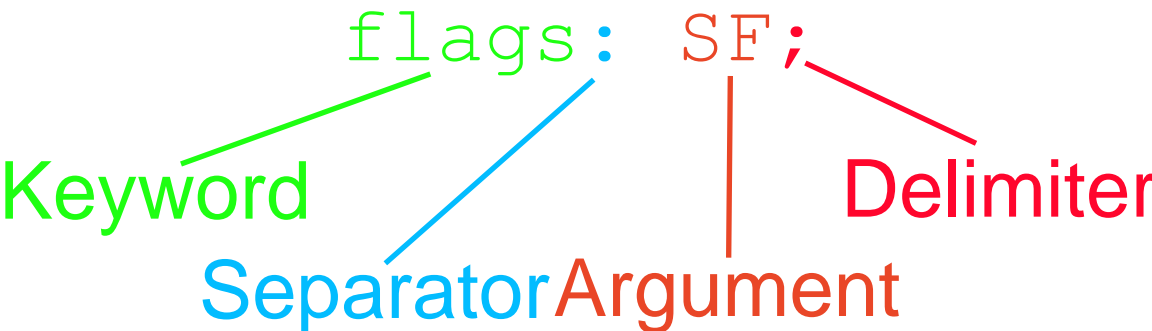
```
(flags: SF; msg: "SYN-FIN scan";)
```



Option start/finish

Option Detail

flags: SF;



Keyword

Separator

Argument

Delimiter

Rule Actions

- The rule action tells Snort what to do when it finds a packet that matches the rule criteria
- Example rule actions:
 - alert: generate an alert and logs the packet
 - log: log the packet
 - pass: ignore the packet
 - activate: alert and turn on another rule
 - dynamic: remains idle until activated by an activate rule, then acts as a log rule

Simple Snort Rules

- e.g., record all TCP traffic inbound for port 79 (finger) going to 10.1.1.0/24

```
log tcp any any -> 10.1.1.0/24 79
```

- e.g., port range and bi-directional traffic

```
log tcp 192.168.1.0/24 0:1023 <> !192.168.1.0/24 any
```

- There is no <- operator

Simple Snort Rules

- You can define variables in rules using var
- Example: detect traffic that goes to CUHK

```
var HOME_NET 192.168.1.0/24
var CUHK_NET 137.189.0.0/16
alert tcp $HOME_NET any -> $CUHK_NET any
```

Snort Rules with Options

- Look for HTTP packets that have substring “/cgi-bin/phf” in payload, and output an alert with message “PHF probe!”

```
alert tcp any any <> any 80 \  
  (content: "/cgi-bin/phf"; msg: "PHF probe!"; sid: 1000001)
```

- You can match a byte pattern:
 - content: “E8C0 FFFF FF|/bin/sh”;
- sid specifies the rule ID
 - >1000000 means user-specific rule ID

Snort Rules with Options


- pcre enables payload inspection with Perl compatible regular expression

```
alert tcp any any -> any any \  
  (pcre: "/^GE*/i"; msg: "HTTP GET"; sid: 1000000)
```

Flow of Snort

➤ Life of a packet inside Snort

- **Decode**: parse a packet (in byte stream) from libpcap and specify the packet structure
- **Preprocess**: advanced decoding, attack detection (will be discussed shortly)
- **Detect**: checks each packet against the various options listed in the Snort rules files (i.e., for rule matching). Each of the keyword options is a plugin.



Preprocessor

- Rule-matching design aims at matching packet data against signature patterns
 - Do packet by packet independently
 - Stateless approach, not trying to correlate two packets
- **Preprocessors** are modules that enable us to do stateful anomaly detection
 - Non-rule based
 - can correlate behavioral patterns of many packets

Preprocessor

- For example, call an existing preprocessor `sfportscan` that detects portscan events
- Create a rule file that has the following:

```
preprocessor sfportscan:  
  proto { all }  
  scan_type { all }  
  sense_level { low }
```


Preprocessor

- Snort provides preprocessors for detecting ARP spoofing

```
preprocessor arpspoof
preprocessor arpspoof_detect_host: 10.0.1.7 00:0C:29:A6:F3:DD
preprocessor arpspoof_detect_host: 10.0.1.10 00:0C:29:AA:BB:CC
```

- Sample output:

```
10/02-11:07:04.360674  [**] [112:4:1] (spp_arpspoof) Attempted ARP cache overwrite attack [**]
10/02-11:07:04.360719  [**] [112:4:1] (spp_arpspoof) Attempted ARP cache overwrite attack [**]
```

...

Write your Own Preprocessor

- Snort implements preprocessors as **plug-in modules**, and allows you to add new preprocessors in a systematic way
- Motivations of adding new preprocessors:
 - try your new research ideas
 - work with existing detection modules


Adding a Preprocessor

- See templates spp_mydetect.*
- Functions a new preprocessor should do:
 - **Setup**: register the preprocessor keyword and initialization function
 - **Init**: parse arguments and link other functions
 - **Main**: process packets and do state tracking
 - **Cleanup**: cleanup when Snort is exiting

Adding a Preprocessor

- How to include spp* into Snort?
 - Insert an include directive in `plugbase.c`
 - Insert a call to your plugin `Setup()` in `RegisterPreprocessors()`
 - Snort keeps a list of preprocessors and executes them one by one, if specified in the config file.
 - Add your plugin code and header file to `Makefile.am` and `Makefile.in` under the `preprocessors/` directory
 - Run `./configure` and make again
- Demo.

Deploying NIDS



- We need to a packet capture device to collect and duplicate packet copies to the NIDS
- Ideally, we shouldn't drop packets
- **Hub:**
 - a half-duplex device that repeats a packet on every output interface
 - a device can see all traffic by enabling promiscuous mode
 - slow capture (hence high packet drops)

R. Bejtlich, "Extrusion Detection", Chapter 4.

Deploying NIDS

➤ Switch:

- Known as port mirroring
- Has a SPAN (switched port analyzer) port that mirrors traffic received on other ports
- Full duplex mode
- Faster than hubs, but may not fast enough for Gigabit networks
- Product example: Cisco Catalyst Switches, Netgear switches



Cisco Catalyst 2960 series switches

Deploying NIDS

➤ Tap / Capture card:

- A device specifically for packet capture
- Has its own packet ring buffer to buffer traffic bursts
- full duplex mode
- Either internal (e.g., PCI cards) or external devices
- Product examples: Endace capture cards, Net Optics taps



DAG 9.2X2 10Gb card

References

➤ Firewall

- Stallings, Ch. 22
- iptables tutorial,
<http://www.netfilter.org/documentation/HOWTO//packet-filtering-HOWTO-7.html>

➤ NIDS

- V. Paxson, “Bro: A System for Detecting Network Intruders in Real-Time”, USENIX 1998.
- SNORT Users Manual 2.8.6, Ch. 1 and 3.
- Brian Caswell, “Harnessing the Power of Snort”